

1) Consider the following relational schemas . . .

- ~~Ans~~
- a) Employee (ename, ecity, state)
 - b) Works (ename, company name, salary)
 - c) Company (company name, city)
 - d) Manager (ename, managername)

Write the expressions to solve the following queries using Relational Algebra, Tuple Relational Calculus and Domain relational Calculus.

a) Write the expression to find the name of the employees getting salary more than 60,000 & working in TCS and WIPRO.

Ans → Relational Algebra → $\Pi[\text{ename}] (\sigma[\text{salary} > 60000 \wedge (\text{company-name} = 'TCS' \vee \text{company-name} = 'WIPRO')] (\text{Works}))$

Tuple Relational Calculus → $\{ w.\text{name} \mid \text{Works}(w) \wedge w.\text{salary} > 60000 \wedge (\text{w.company-name} = 'TCS' \vee \text{w.company-name} = 'WIPRO') \}$

Domain Relational Calculus → $\{ \text{ename} \mid \exists \text{company-name}, \text{salary} (\text{Works}(\text{ename}, \text{company-name}, \text{salary}) \wedge \text{salary} > 60000 \wedge (\text{company-name} = 'TCS' \vee \text{company-name} = 'WIPRO')) \}$

b) Write the expressions to find the salary and name of the employees working in any company placed in DELHI.

Ans → Relational Algebra → $\Pi[\text{ename}, \text{salary}] (\sigma[\text{city} = 'DELHI'] (\text{Company} \bowtie \text{Company} \cdot \text{company-name} = \text{Works} \cdot \text{company-name} \text{ Works}))$

Tuple Relational Calculus (TRC) → $\{ \langle w.\text{ename}, w.\text{salary} \rangle \mid \text{Work}(w) \wedge \text{Company}(c) \wedge w.\text{company-name} = c.\text{company-name} \wedge c.\text{city} = 'DELHI' \}$

Domain Relational Calculus (DRC) → $\{ \langle \text{ename}, \text{salary} \rangle \mid \exists \text{company-name}, \text{city} (\text{Works}(\text{ename}, \text{company-name}, \text{salary}) \wedge \text{Company}(\text{company-name}, \text{city}) \wedge \text{city} = 'DELHI') \}$

c) Write the expression to find the salary of the employees living in city MUMBAI and working under the manager JOHN.

Ans → Relational Algebra → $\Pi[\text{salary}] (\sigma[\text{ecity} = 'MUMBAI'] (\text{Employee} \bowtie \text{Employee.ename} = \text{Works.ename} \text{ Works} \bowtie \text{Works.ename} = \text{Manager.ename} \bowtie [\text{managername} = 'JOHN'] (\text{Manager})))$

Tuple Relational Calculus $\rightarrow \{ \langle w \cdot \text{salary} \rangle \mid \text{Works}(w) \wedge \text{Employee}(e) \wedge \text{Manages}(m) \wedge$

$w \cdot \text{ename} = e \cdot \text{ename} \wedge w \cdot \text{ename} \wedge e \cdot \text{ecity} = \text{'MUMBAI'} \wedge m \cdot \text{managername} = \text{'JOHN'}$ }

Domain Relational Calculus $\rightarrow \{ \langle \text{salary} \rangle \mid \exists \text{ename, ecity, state, company-name, managername} ($

$\text{Employee}(\text{ename, ecity, state}) \wedge \text{ecity} = \text{'MUMBAI'} \wedge$

$\text{Works}(\text{ename, company-name, salary}) \wedge$

$\text{Manages}(\text{ename, managername}) \wedge \text{managername} = \text{'JOHN'}) \}$

Q2) Consider the following three relations in a relational database: Employee(eId, name), Brand(bId, bName), Own(eId, bId). Write the relational algebra expression that will return the set of eIds who own all the brands. Write the equivalent tuple relational calculus & domain relational calculus query.

Ans \rightarrow Relational Algebra \rightarrow Let take $R = \prod_{eId, bId} \text{Own}$, $S = \prod_{bId} \text{Brand}$

$\prod_{[eId]}(\text{Own}) - \prod_{[eId]}(\prod_{[eId, bId]}(\text{Own}) \times \prod_{[bId]}(\text{Brand})) - \text{Own}$

Tuple Relational Calculus $\rightarrow \{ e \cdot eId \mid \text{Employee}(e) \wedge \forall b (\text{Brand}(b) \rightarrow \exists o (\text{Own}(o) \wedge o \cdot eId$

$= e \cdot eId \wedge o \cdot bId = b \cdot Id)) \}$

Domain Relational Calculus $\rightarrow \{ \langle eId \rangle \mid \exists \text{name} (\text{Employee}(eId, name)) \wedge \forall bId, bName$

$(\text{Brand}(bId, bName)) \rightarrow \exists eId2 (\text{Own}(eId2, bId) \wedge eId2 = eId) \}$

Q3) Check whether the given schedules are conflict serializable or not-

a) S1: r1(x) r1(y) r2(x) r2(y) w2(y) w1(x)

b) S2: r1(x) r2(x) r2(y) w2(y) r1(y) w1(x)

Let $r_i(z)$ and $w_i(z)$ denote read and write operations respectively on a data item z by a transaction T_i . Consider the following two schedules.

Ans \rightarrow a) Schedule S1 \rightarrow r1(x) r1(y) r2(x) r2(y) w2(y) w1(x) Transactions $\rightarrow T_1 \& T_2$

Conflicts \rightarrow

i) r1(x) & w1(x) \rightarrow same $T_1 \rightarrow$ no conflict

ii) r2(x) & r1(x) \rightarrow read-read \rightarrow no conflict

iii) r2(x) & w1(x) \rightarrow r2(x) before w1(x) \rightarrow conflict $\rightarrow T_2 \rightarrow T_1$

iv) r1(y) & w2(y) \rightarrow T1 reads before T2 writes \rightarrow conflict $\rightarrow T_1 \rightarrow T_2$

v) r2(y) & w2(y) \rightarrow same $T_2 \rightarrow$ no conflict

Precedence Graph Edges $\rightarrow T_1 \rightarrow T_2$ (because of $rd(y)$ before $w_2(y)$)
 $T_2 \rightarrow T_1$ (because of $rd(x)$ before $w_1(x)$)

Cycle detected $\rightarrow T_1 \rightarrow T_2 \rightarrow T_1$

$\therefore S_1$ is not conflict serializable.

b) Schedule $S_2 \rightarrow S_2$: $rd(x) \& rd(y) \& rd(y) \& w_2(y) \& rd(y) \& w_1(x)$

- Conflict \rightarrow
- i) $rd(x) \& rd(x) \rightarrow$ read-read \rightarrow no conflict
 - ii) $rd(y) \& rd(y) \rightarrow$ both read \rightarrow no conflict
 - iii) $w_2(y) \& rd(y) \rightarrow w_2(y)$ before $rd(y) \rightarrow$ conflict $\rightarrow T_2 \rightarrow T_1$
 - iv) $rd(x) \& w_1(x) \rightarrow$ same $T_1 \rightarrow$ no conflict
 - v) $rd(x) \& w_1(x) \rightarrow T_2$ reads before T_1 writes \rightarrow conflict $\rightarrow T_2 \rightarrow T_1$

Precedence Graph Edges $\rightarrow T_2 \rightarrow T_1$ (from both $w_2(y)$ before $rd(y)$ and $rd(x)$ before $w_1(x)$)

No edge from T_1 to $T_2 \rightarrow$ no cycle

$\therefore S_2$ is conflict serializable.

4) Check whether the given schedule below is conflict serializable or not justify your answer & show the precedence graph.

T_1	T_5
$read(A)$	
$A := A - 50$	
$write(A)$	
	$read(B)$
	$B := B - 10$
	$write(B)$
$read(B)$	
$B := B + 50$	
$write(B)$	
	$read(A)$
	$A := A + 10$
	$write(A)$

Ans \rightarrow Conflict Detection \rightarrow

- i) On $B \rightarrow$ a) $T_5 : write(B) \rightarrow T_1 : read(B) = W-R$ conflict $\rightarrow T_5 \rightarrow T_1$
- b) $T_5 : write(B) \rightarrow T_1 : write(B) = W-W$ conflict $\rightarrow T_5 \rightarrow T_1$

Contd ii) On A → a) T1: write(A) → T5: read(A) = W-R conflict $\rightarrow T1 \rightarrow T5$
b) T1: write(A) → T5: write(A) = W-W conflict $\rightarrow T1 \rightarrow T5$

Precedence Graph \rightarrow Nodes $\rightarrow T1, T5$

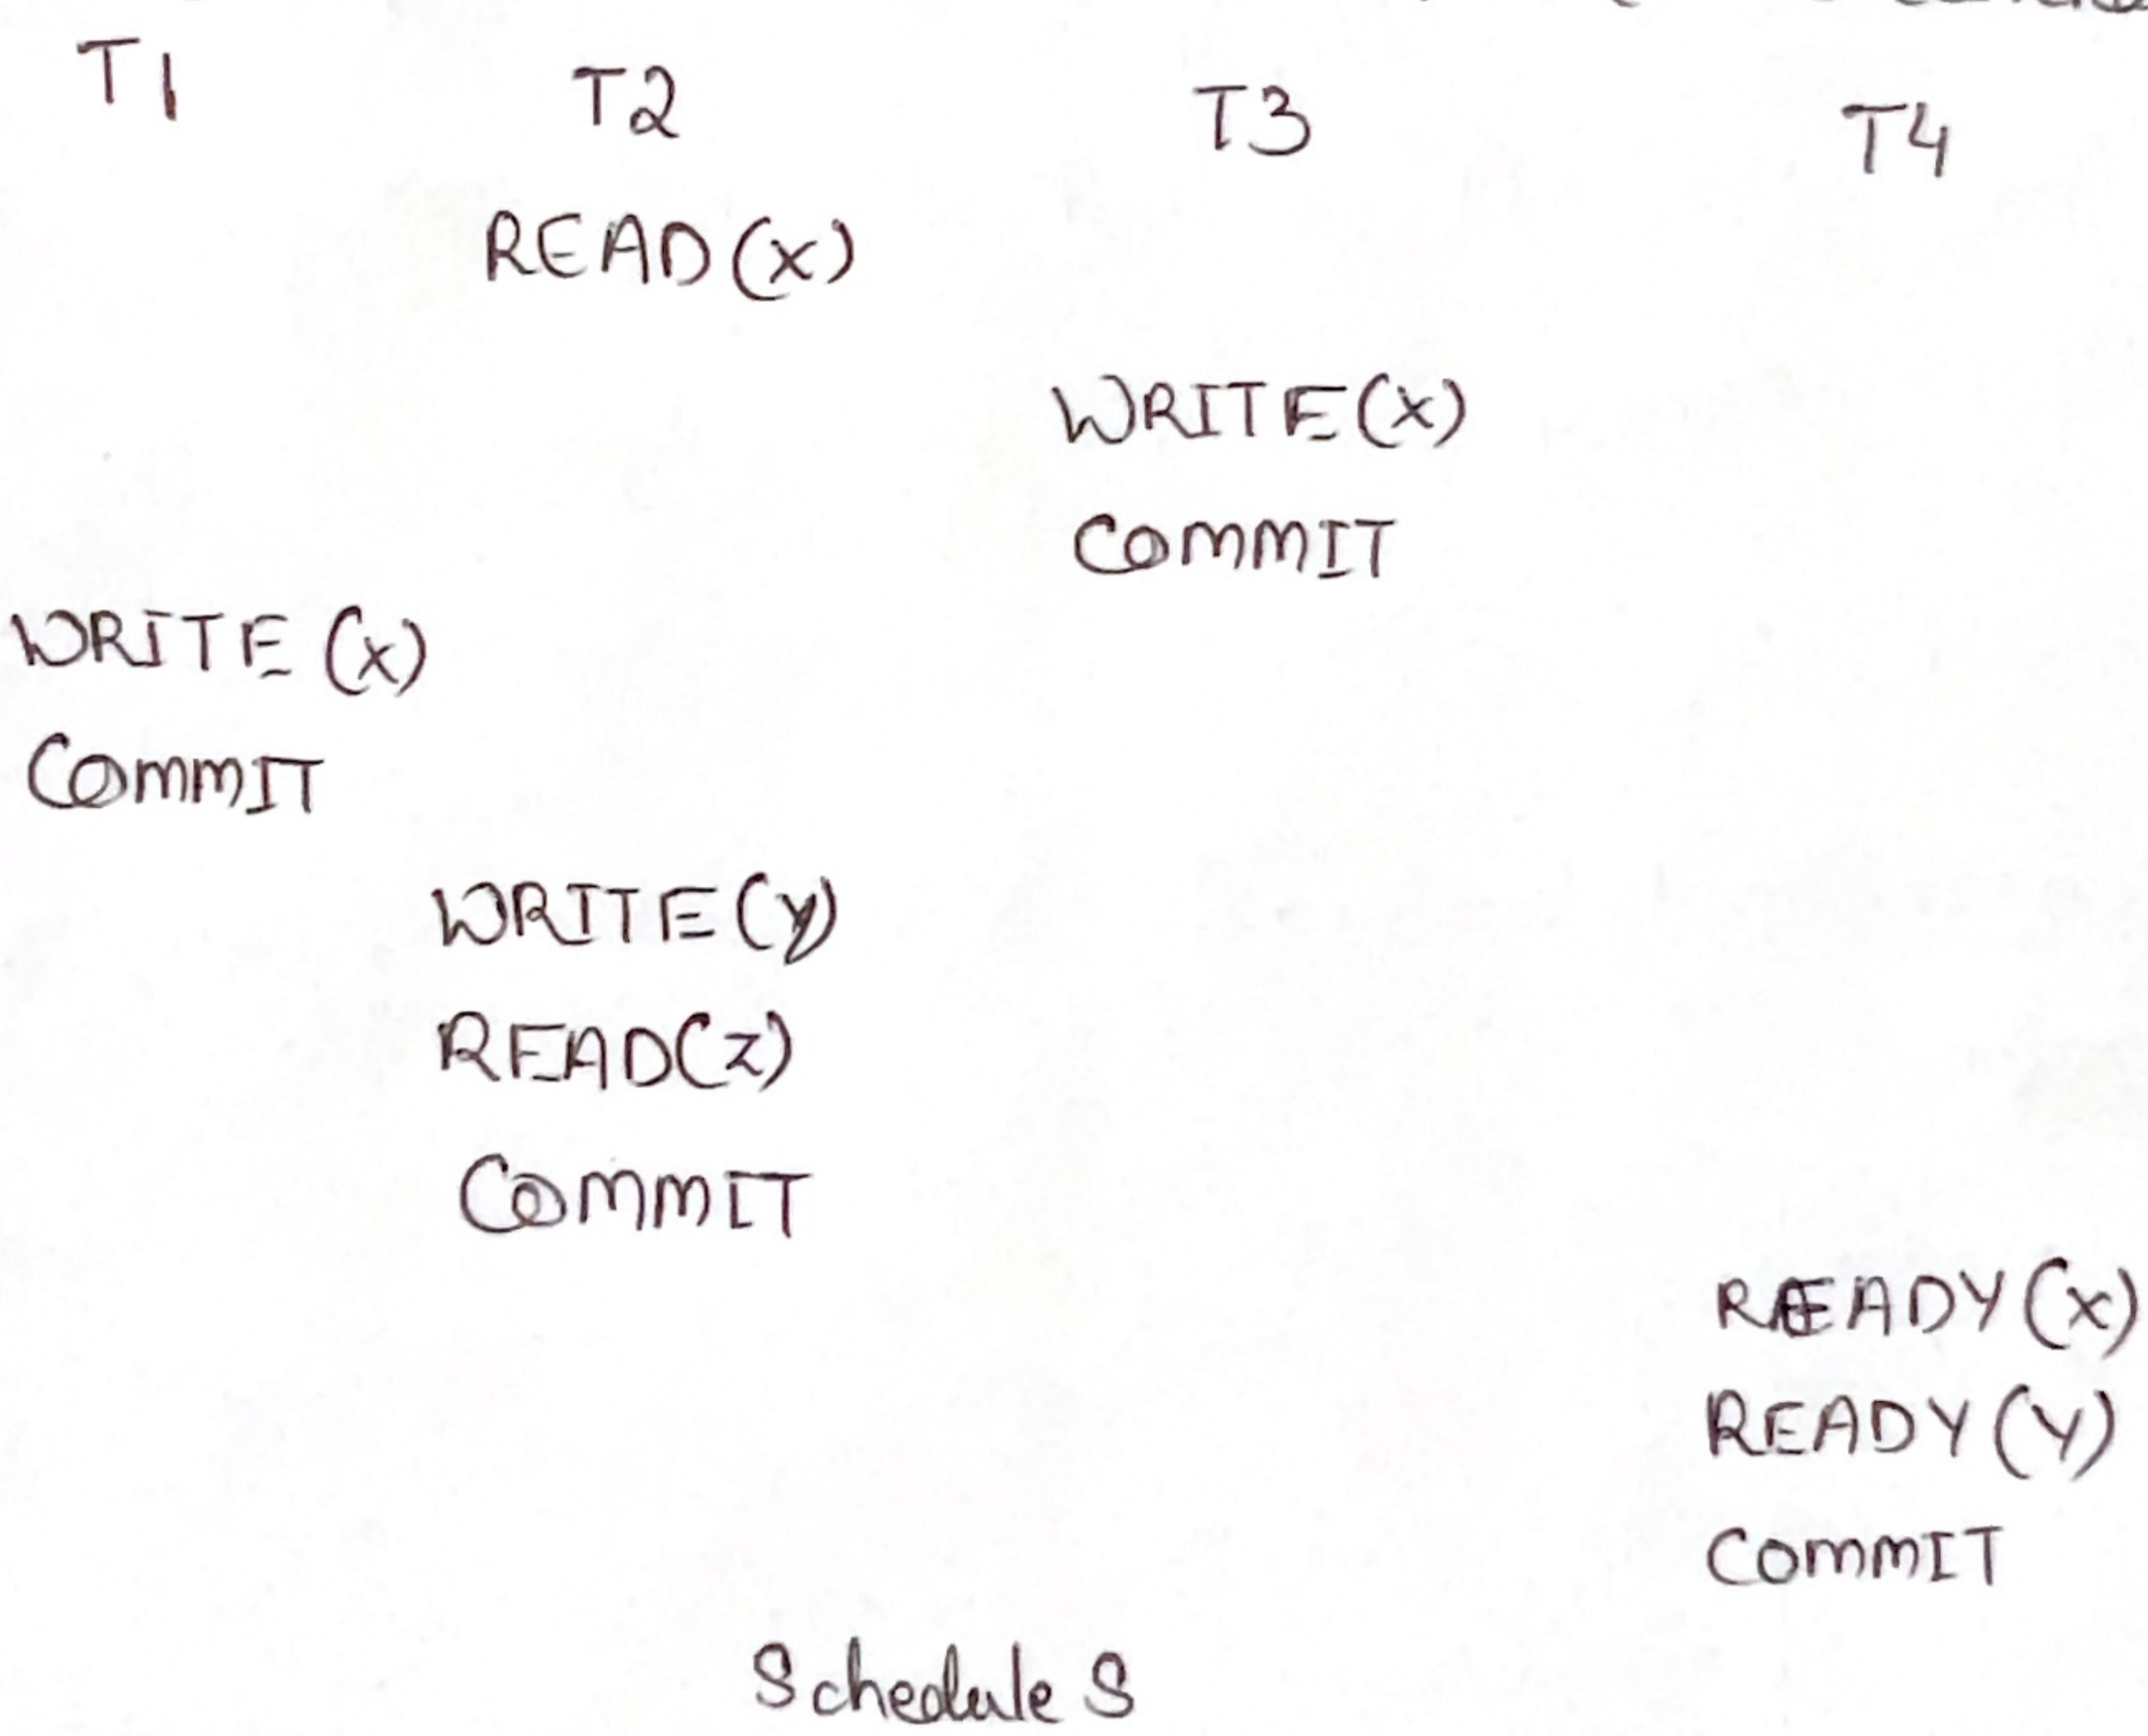
Edges $\rightarrow T1 \rightarrow T5$ (due to conflict on A)

$T5 \rightarrow T1$ (due to conflict on B)

Cycle detected $\rightarrow T1 \rightarrow T5 \rightarrow T1$

Conclusion \rightarrow The schedule is not conflict serializable because the precedence graph contains a cycle.

5) Check whether the given schedule S is conflict serializable and recoverable or not.



Ans \rightarrow Conflicts on X \rightarrow

- i) T2: READ(x) before T3: WRITE(x) $\rightarrow T2 \rightarrow T3$
- ii) T3: WRITE(x) before T1: WRITE(x) $\rightarrow T3 \rightarrow T1$
- iii) T1: WRITE(x) before T4: READ(x) $\rightarrow T1 \rightarrow T4$

Conflicts on Y \rightarrow T2: WRITE(y) before T4: READ(y) $\rightarrow T2 \rightarrow T4$

Build Precedence Graph \rightarrow Nodes $\rightarrow T1, T2, T3, T4$

Edges \rightarrow $T2 \rightarrow T3$ $T2 \rightarrow T4$
 $T3 \rightarrow T1$
 $T1 \rightarrow T4$

The graph has no cycles.

Conflict Serializable \rightarrow Because the precedence graph is acyclic, the schedule is conflict serializable.

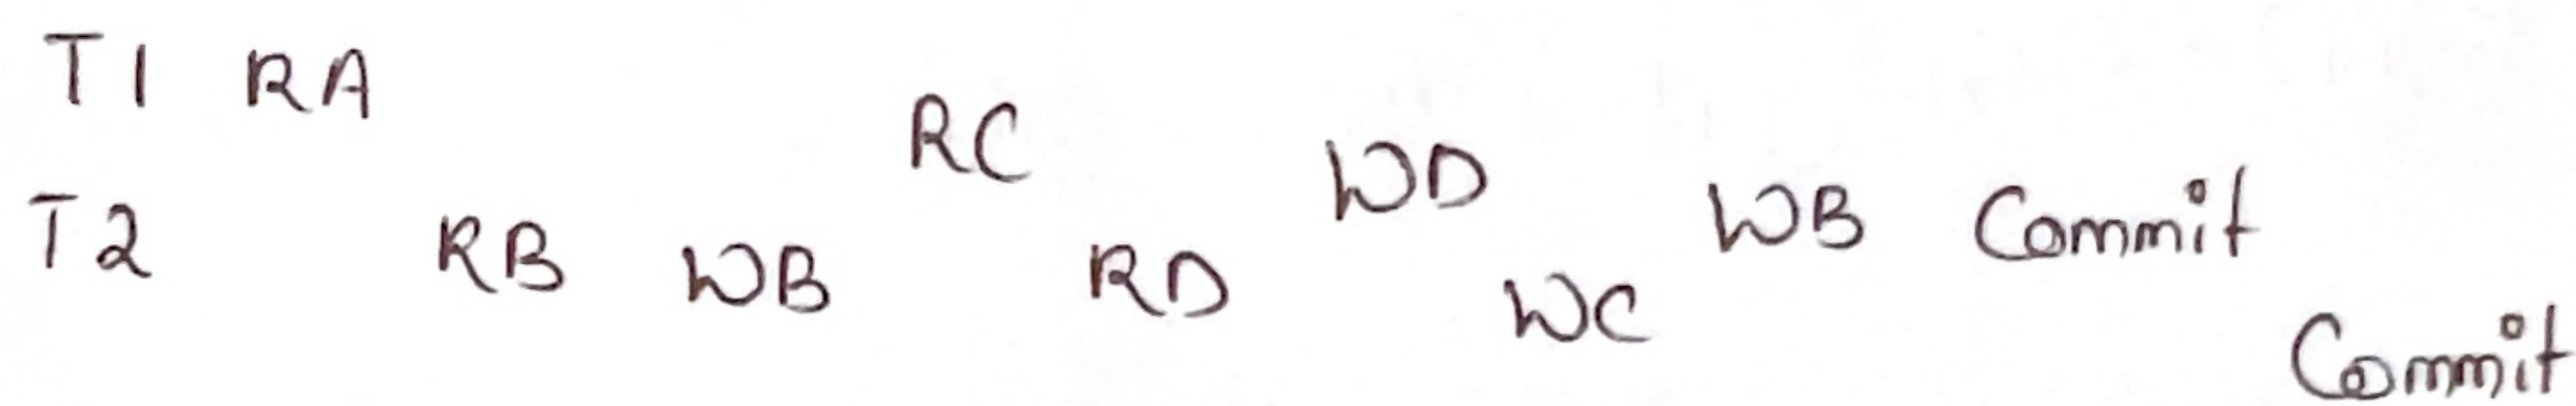
Possible serial order $\rightarrow T_2 \rightarrow T_3 \rightarrow T_1 \rightarrow T_4$

Recoverability \rightarrow A schedule is recoverable if a transaction only commits after the transactions it reads from have committed.

- T_2 reads X before T_3 writes it \rightarrow No read-after-write conflict
- T_4 reads X written by T_1 , and Y written by $T_2 \rightarrow T_4$ commits after both T_1 & T_2

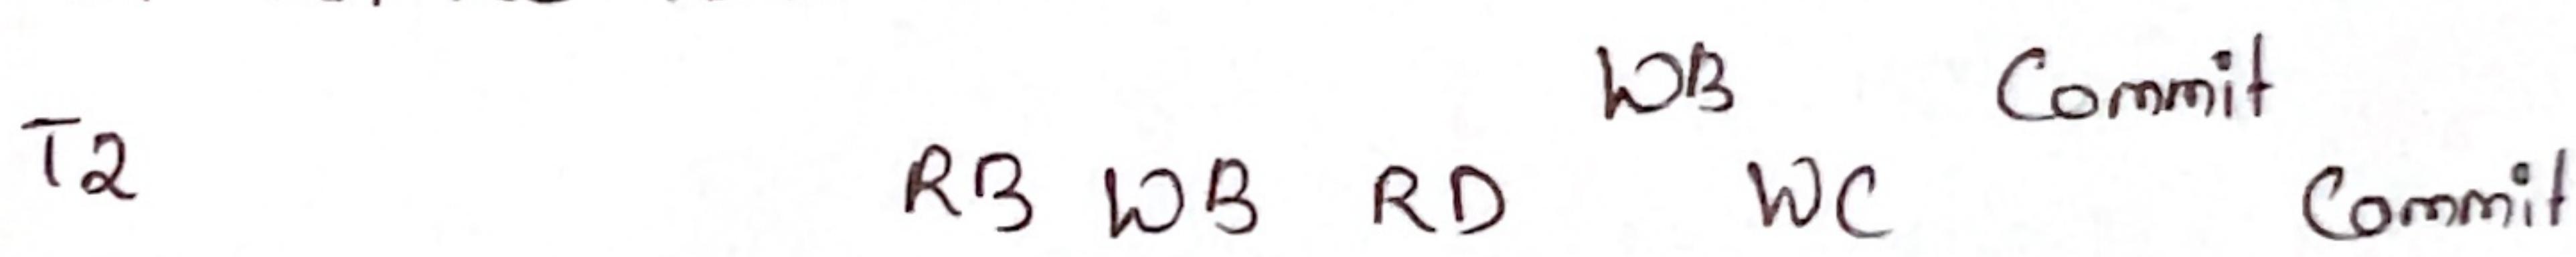
Yes, its recoverable

6) Consider a schedule of Transaction T_1 and T_2 :

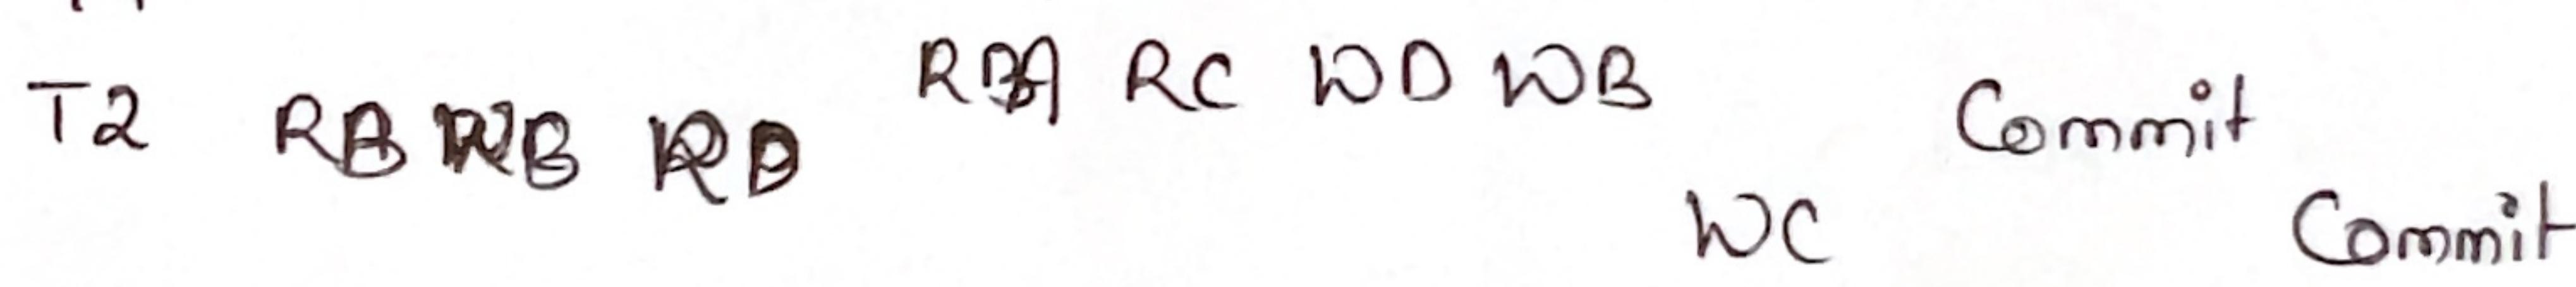


Hence, RX stands for "Read(X)" and WX stands for "Write(X)". Which one of the following schedules is conflict equivalent to the above schedule?

a) ~~RD~~ T1 RA RC WD



b) T1



Ans \rightarrow From the given schedule, Conflicts \rightarrow

- Read-Write (RW) on the same item
- Write-Read (WR) on same item
- Write-Write (WW) on same item.

1) $T_1: WD$ and $T_2: RD \rightarrow$ WR conflict on D $\rightarrow T_1 \rightarrow T_2$

2) $T_1: WB$ and $T_2: WB \rightarrow$ WW conflict on B $\rightarrow T_1 \rightarrow T_2$

3) $T_2: WB$ and $T_1: WB \rightarrow$ WW conflict on B $\rightarrow T_2 \rightarrow T_1$

So, this gives a cycle $\rightarrow T_1 \rightarrow T_2 \rightarrow T_1 \rightarrow$ Not conflict serializable

\therefore The original schedule is not conflict serializable.

In option \rightarrow

a) Same as the original schedule \rightarrow Also not conflict serializable.

b) Conflicts \rightarrow i) $T_2: WB$ before $T_1: WB \rightarrow$ WW conflict $\rightarrow T_2 \rightarrow T_1$

ii) $T_1: WD$ before $T_2: RD \rightarrow$ WR conflict $\rightarrow T_1 \rightarrow T_2$

$T_1 \rightarrow T_2$ and $T_2 \rightarrow T_1 \Rightarrow$ Cycle exists, not serializable

So, Neither (A) nor (B) is conflict equivalent to a conflict-serializable schedule.

7) Consider the following two transactions \rightarrow

$T_1 : \text{read}(A);$
 $\text{read}(B);$
 $\text{if } A=0 \text{ then } B := B+1;$
 $\text{write}(B).$

$T_2 : \text{read}(B);$
 $\text{read}(A);$
 $\text{if } B=0 \text{ then } A := A+1;$
 $\text{write}(A).$

Add lock and unlock instructions to transactions T_1 & T_2 , so that they observe the two-phase locking protocol. Can the execution of these transactions result in a deadlock?

Ans \rightarrow Apply Two-Phase Locking (2PL) \rightarrow • Growing phase: Acquire all locks (shared or exclusive) - no unlocks.

• Shrinking phase: Release locks - no new locks allowed.

We assume that: • Shared Lock(S) for reading.
• Exclusive Lock(X) for writing.

T_1 with locks \rightarrow

LOCK-S(A)
 $\text{read}(A)$
LOCK-S(B)
 $\text{read}(B)$
UNLOCK-S(A)
 $\text{if } A=0 \text{ then}$
LOCK-X(B)
 $B := B+1$
 $\text{write}(B)$
UNLOCK-X(B)
UNLOCK-S(B)

T_2 with locks \rightarrow

LOCK-S(B)
LOCK-S(A)
LOCK-X(A)
 $\text{read}(B)$
 $\text{read}(A)$
 $\text{if } B=0 \text{ then}$
 $A := A+1$
 $\text{write}(A)$
UNLOCK-X(A)
UNLOCK-X(B)
UNLOCK-X(C)

Yes - a classic deadlock situation is possible.

Suppose:

- T_1 locks A (shared), then requests lock on B
- T_2 locks B (shared), then request lock on A

Then:

• T_1 is waiting for B (held by T_2)

• T_2 is waiting for A (held by T_1)

\therefore This is a circular wait \Rightarrow Deadlock

8) Check whether the given schedule below can complete all its operations in the given order using timestamp based protocol or not.

~~Ans~~ Ans → Time Stamp Protocol rules →

For a data item, x :

Read rule:

- If $TS(T) < W_TS(x) \rightarrow$ Abort T (younger T is trying to read outdated value)
- Else → allow, and set $R_TS(x) = \max(R_TS(x), TS(T))$

Write rule:

- If $TS(T) < R_TS(x) \rightarrow$ Abort T
- If $TS(T) < W_TS(x) \rightarrow$ Abort T
- Else → allow, and update $W_TS(x)$

Let's denote: Initial: $R_TS(A) = 0, W_TS(A) = 0$

Step 1: T1 reads A

- Allowed since $TS(T1) < W_TS(A)$ is false ($T1$ is older)
- Update $R_TS(A) = TS(T1)$

Step 2: T2 reads A

- Allowed if $TS(T2) \geq W_TS(A) \rightarrow$ true ($W_TS(A) = 0$)
- Update $R_TS(A) = TS(T2)$;

Step 3: T1 writes A

- Allowed if: $TS(T1) \geq R_TS(A)$ may be false now ($TS(T1) < TS(T2) = R_TS(A)$) →

So T1 will fail to write A.

Abort T1

Step 4: T2 writes A

- T2 is newer, $R_TS(A) = TS(T2)$, so
 - $TS(T2) \geq R_TS(A)$
 - $TS(T2) \geq W_TS(A)$
 - → Allowed

∴ No, the schedule cannot complete all operations using timestamp-based protocol.

9) Consider T_0 & T_1 be two transactions. T_0 transfers 50Rs from account A to account B and T_1 withdraws 100Rs from account C. Both transaction can't commit successfully. Write down the system log corresponding to T_0 & T_1 to handle any failures.

Ans → Initial bal $\rightarrow A = 1000, B = 2000, C = 700$

Transactions $\rightarrow T_0$: Transfer ₦50 from A to B ; T_1 : Withdraw ₦100 from C

$$\rightarrow A = A - 50 \rightarrow 950$$

$$\rightarrow C = C - 100 \rightarrow 600$$

$$\rightarrow B = B + 50 \rightarrow 2050$$

System Log \rightarrow START T_0

- $\langle T_0, X, \text{old-value}, \text{new-value} \rangle$

- Commit T_0

Log Entries \rightarrow START T_0

$\langle T_0, A, 1000, 950 \rangle$

$\langle T_0, B, 2000, 2050 \rangle$

COMMIT T_0

START T_1

$\langle T_1, C, 700, 600 \rangle$

COMMIT T_1

10) Write down Two major differences between immediate and deferred database modification.

Ans → i) When the changes are applied to the database:

- Immediate Modification \rightarrow Changes to the database are applied as soon as the write operation is issued, even before the transaction commits.

- Deferred Modification \rightarrow Changes are not applied immediately; they are deferred until the transaction successfully commits.

ii) Logging Requirements \rightarrow

- Immediate Modification \rightarrow Requires both undo & redo information in the log (since changes happen before commit and need might need to be undone if a crash occurs).

- Deferred Modification \rightarrow Requires only redo information (since changes are applied only after a successful commit, no undo is necessary).